

# Logistic Regression

May 23, 2022

## 1 Imports

```
[1]: from sklearn.model_selection import train_test_split
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

## 2 Load the data

```
[2]: heart_data = pd.read_csv('heart.csv')
heart_data.head()
```

```
[2]:    age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope \
0    63    1    3      145   233    1        0     150      0       2.3      0
1    37    1    2      130   250    0        1     187      0       3.5      0
2    41    0    1      130   204    0        0     172      0       1.4      2
3    56    1    1      120   236    0        1     178      0       0.8      2
4    57    0    0      120   354    0        1     163      1       0.6      2

      ca  thal  target
0      0    1       1
1      0    2       1
2      0    2       1
3      0    2       1
4      0    2       1
```

## 3 Data exploration

```
[3]: heart_data.shape
```

```
[3]: (303, 14)
```

```
[4]: heart_data.describe()
```

```
[4]:
```

	age	sex	cp	trestbps	chol	fbns	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	

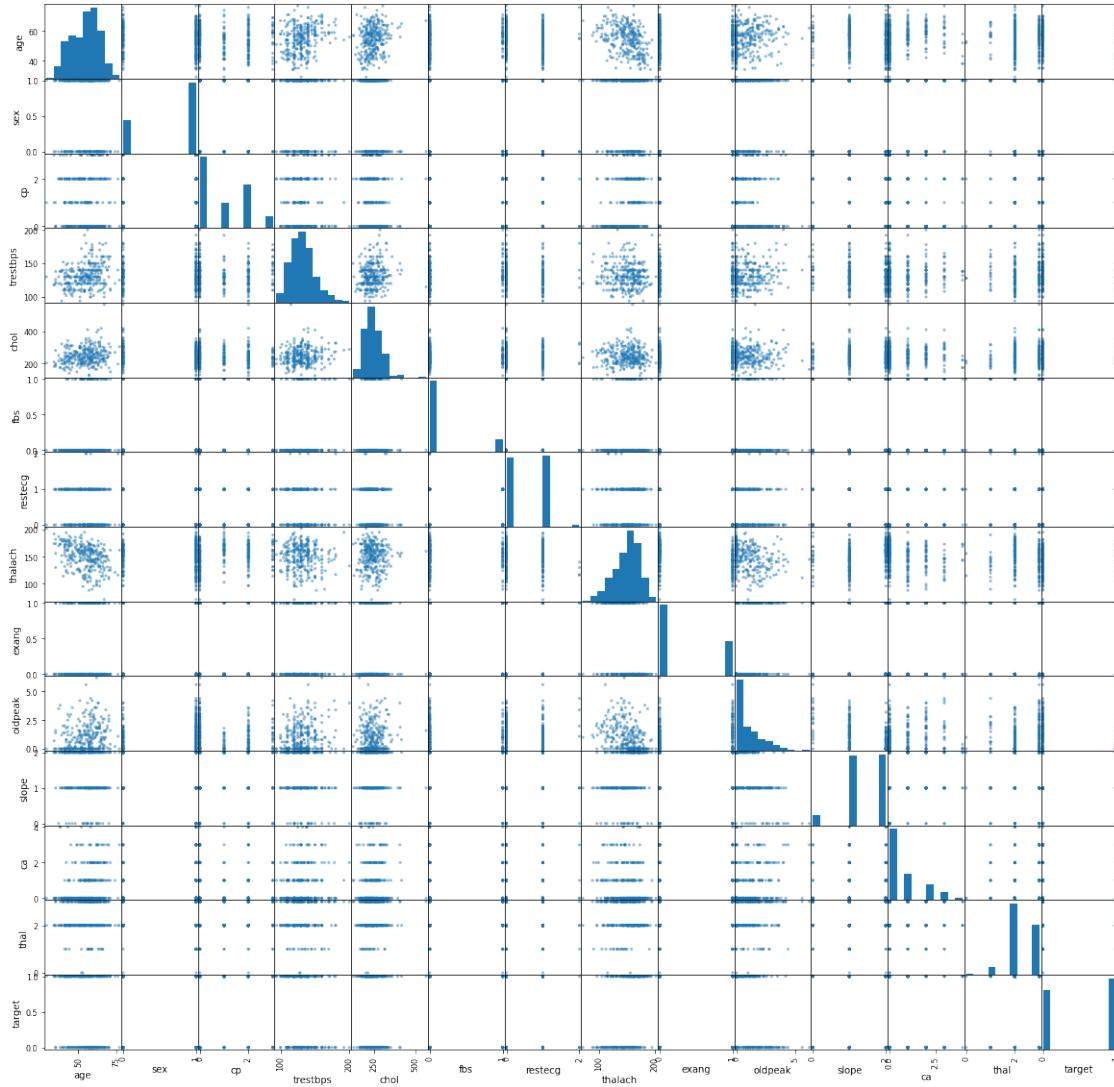
  

	restecg	thalach	exang	oldpeak	slope	ca	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	
mean	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373	
std	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606	
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000	
50%	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000	
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000	
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000	

	thal	target
count	303.000000	303.000000
mean	2.313531	0.544554
std	0.612277	0.498835
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

```
[5]: pd.plotting.scatter_matrix(heart_data, figsize=(20, 20))
plt.savefig('scatter_matrix.jpg')
plt.show()
```



```
[6]: heart_data.corr()
```

```
[6]:      age      sex      cp  trestbps      chol      fbs  \
age    1.000000 -0.098447 -0.068653  0.279351  0.213678  0.121308
sex   -0.098447  1.000000 -0.049353 -0.056769 -0.197912  0.045032
cp     -0.068653 -0.049353  1.000000  0.047608 -0.076904  0.094444
trestbps  0.279351 -0.056769  0.047608  1.000000  0.123174  0.177531
chol    0.213678 -0.197912 -0.076904  0.123174  1.000000  0.013294
fbs     0.121308  0.045032  0.094444  0.177531  0.013294  1.000000
restecg -0.116211 -0.058196  0.044421 -0.114103 -0.151040 -0.084189
thalach -0.398522 -0.044020  0.295762 -0.046698 -0.009940 -0.008567
exang   0.096801  0.141664 -0.394280  0.067616  0.067023  0.025665
oldpeak  0.210013  0.096093 -0.149230  0.193216  0.053952  0.005747
slope   -0.168814 -0.030711  0.119717 -0.121475 -0.004038 -0.059894
```

```

ca          0.276326  0.118261 -0.181053  0.101389  0.070511  0.137979
thal        0.068001  0.210041 -0.161736  0.062210  0.098803 -0.032019
target     -0.225439 -0.280937  0.433798 -0.144931 -0.085239 -0.028046

```

```

            restecg    thalach     exang   oldpeak      slope       ca \
age      -0.116211 -0.398522  0.096801  0.210013 -0.168814  0.276326
sex      -0.058196 -0.044020  0.141664  0.096093 -0.030711  0.118261
cp       0.044421  0.295762 -0.394280 -0.149230  0.119717 -0.181053
trestbps -0.114103 -0.046698  0.067616  0.193216 -0.121475  0.101389
chol     -0.151040 -0.009940  0.067023  0.053952 -0.004038  0.070511
fbs      -0.084189 -0.008567  0.025665  0.005747 -0.059894  0.137979
restecg   1.000000  0.044123 -0.070733 -0.058770  0.093045 -0.072042
thalach   0.044123  1.000000 -0.378812 -0.344187  0.386784 -0.213177
exang     -0.070733 -0.378812  1.000000  0.288223 -0.257748  0.115739
oldpeak   -0.058770 -0.344187  0.288223  1.000000 -0.577537  0.222682
slope     0.093045  0.386784 -0.257748 -0.577537  1.000000 -0.080155
ca        -0.072042 -0.213177  0.115739  0.222682 -0.080155  1.000000
thal     -0.011981 -0.096439  0.206754  0.210244 -0.104764  0.151832
target    0.137230  0.421741 -0.436757 -0.430696  0.345877 -0.391724

```

```

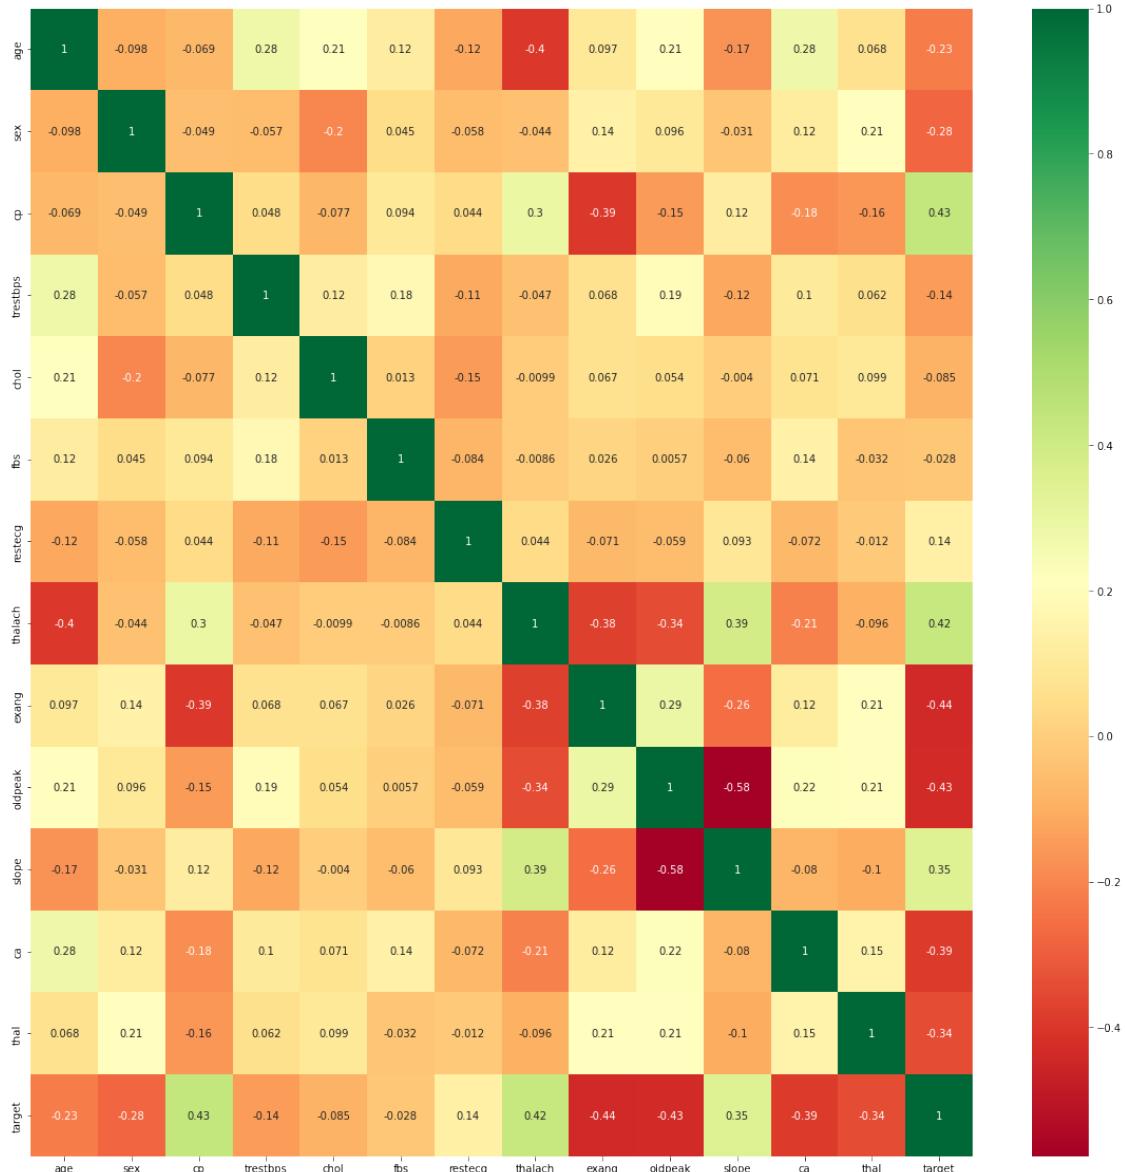
            thal      target
age      0.068001 -0.225439
sex      0.210041 -0.280937
cp       -0.161736  0.433798
trestbps 0.062210 -0.144931
chol     0.098803 -0.085239
fbs      -0.032019 -0.028046
restecg -0.011981  0.137230
thalach -0.096439  0.421741
exang    0.206754 -0.436757
oldpeak  0.210244 -0.430696
slope    -0.104764  0.345877
ca       0.151832 -0.391724
thal     1.000000 -0.344029
target   -0.344029  1.000000

```

```

[7]: plt.figure(figsize=(20,20))
sns.heatmap(heart_data.corr(), annot=True, cmap="RdYlGn")
plt.savefig('corr.jpg')
plt.show()

```



## 4 Logistic Regression

```
[8]: class LogiReg:
    def __init__(self, lr=0.001, iters_num=100, log=True):
        self.lr = lr
        self.iters_num = iters_num
        self.weights = None
        self.log = log
```

```

def sigmoid(self, x):
    return 1 / (1 + np.exp(-x))

def b_intercept(self, X):
    intercept = np.ones((X.shape[0], 1))
    # concatenate them to the value of X
    return np.concatenate((intercept, X), axis=1)

def fit (self, X, y):
    # initialize parameters
    X = self.b_intercept(X)
    self.weights = np.zeros(X.shape[1])

    # gradient descent
    for i in range(self.iters_num):

        z = np.dot(X, self.weights)
        # apply sigmoid function
        y_pred = self.sigmoid(z)

        # compute gradients
        grad = np.dot(X.T, (y_pred - y)) / y.size
        self.weights -= self.lr * grad

        # new W * Xi
        new_z = np.dot(X, self.weights)
        y_pred = self.sigmoid(new_z)

        # calculate loss
        loss = self.loss(y_pred, y)
        if i % 100 == 0 and self.log:
            print(f'iteration #{i}, loss = {loss}')

def loss(self, y_pred, y):
    return (-y * np.log(y_pred) - (1 - y) * np.log(1 - y_pred)).mean()

# predict the probability values
def predict_prob(self, X):
    X = self.b_intercept(X)
    z = np.dot(X, self.weights)
    return self.sigmoid(z)

# predict the actual values 0 or 1 using round
def predict(self, X):

```

```
    return self.predict_prob(X).round()
```

```
[9]: model = LogiReg(0.01, 3000)
```

## 5 Data preprocessing

```
[10]: # creating features and labels matrices
X = heart_data.drop(['target'], axis=1) # axis =1 because we deal with
    ↵column
y = heart_data['target']
```

```
[11]: # normalization step
X = (X-X.mean())/X.std()
```

```
[12]: # train test splitting
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↵random_state=0)
```

## 6 All features

### 6.0.1 Model training

```
[13]: model.fit(X_train, y_train)
```

```
iteration #0, loss = 0.6899513425649656
iteration #100, loss = 0.5080874138342651
iteration #200, loss = 0.4434654806586477
iteration #300, loss = 0.4126782301146911
iteration #400, loss = 0.395135628962435
iteration #500, loss = 0.3839824082850972
iteration #600, loss = 0.37636232794764624
iteration #700, loss = 0.37088798112082094
iteration #800, loss = 0.3668077919367627
iteration #900, loss = 0.36368039187628926
iteration #1000, loss = 0.36123006516245904
iteration #1100, loss = 0.35927600560309153
iteration #1200, loss = 0.35769492735406383
iteration #1300, loss = 0.3564000366595718
iteration #1400, loss = 0.35532858424449393
iteration #1500, loss = 0.354434174928488
iteration #1600, loss = 0.35368184113103557
iteration #1700, loss = 0.3530447876229182
iteration #1800, loss = 0.3525021822418355
iteration #1900, loss = 0.3520376211278857
iteration #2000, loss = 0.3516380404969885
iteration #2100, loss = 0.3512929309376693
```

```

iteration #2200, loss = 0.3509937609081084
iteration #2300, loss = 0.3507335475692134
iteration #2400, loss = 0.3505065331009784
iteration #2500, loss = 0.3503079376649882
iteration #2600, loss = 0.3501337688128113
iteration #2700, loss = 0.3499806729763713
iteration #2800, loss = 0.34984581868595993
iteration #2900, loss = 0.34972680395795647

```

## 6.0.2 Evaluation

```
[14]: def accuracy(y, y_pred):
    accuracy = np.mean(y == y_pred)
    return accuracy

def plot_regression(X, y, model):
    plt.figure(figsize=(10, 6))
    plt.scatter(X[y == 0][:, 0], X[y == 0][:, 1], color='b', label='0')
    plt.scatter(X[y == 1][:, 0], X[y == 1][:, 1], color='r', label='1')
    plt.legend()
    x1_min, x1_max = X[:,0].min(), X[:,0].max(),
    x2_min, x2_max = X[:,1].min(), X[:,1].max(),
    xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max), np.linspace(x2_min, x2_max))
    grid = np.c_[xx1.ravel(), xx2.ravel()]
    probs = model.predict_prob(grid).reshape(xx1.shape)
    plt.contour(xx1, xx2, probs, [0.5], linewidths=1, colors='black')
    plt.show()
```

```
[15]: predictions = model.predict(X_test)
print("Accuracy:", accuracy(y_test, predictions))
```

Accuracy: 0.8360655737704918

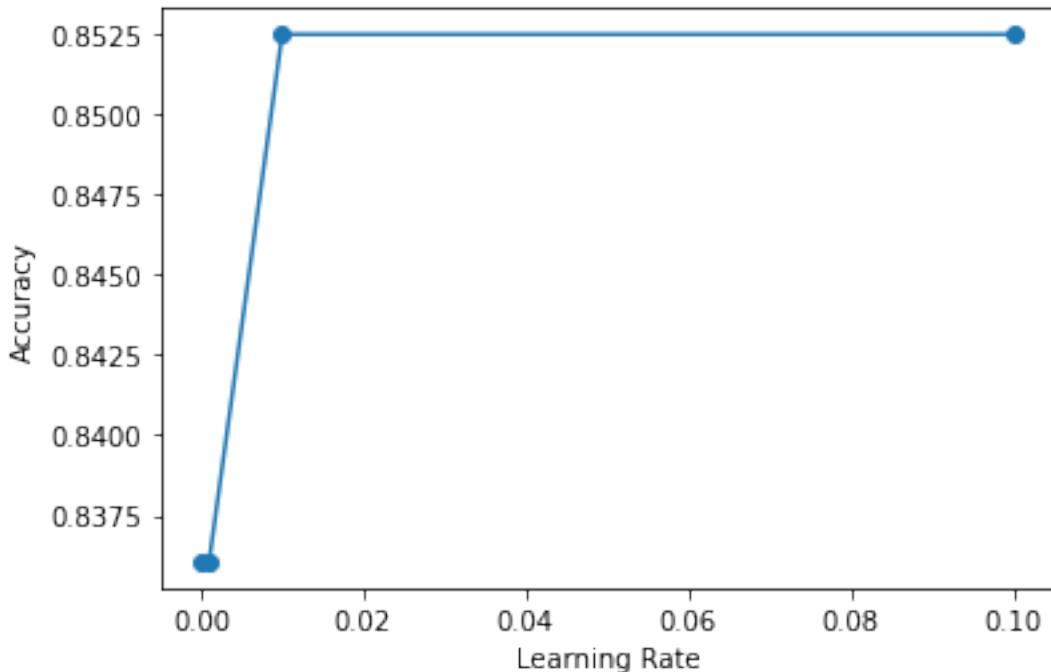
## 7 Learning rate

```
[16]: lrs = [0.1, 0.01, 0.001, 0.0001]
acc = list()
for lr in lrs:
    model = LogiReg(lr, 10000, False)
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    acc.append(accuracy(y_test, predictions))

best_lr = lrs[acc.index(max(acc))]
print('Best learning rate is', best_lr)
```

Best learning rate is 0.1

```
[17]: plt.plot(lrs, acc, '-o')
plt.xlabel('Learning Rate')
plt.ylabel('Accuracy')
plt.savefig('lrs.jpg')
plt.show()
```



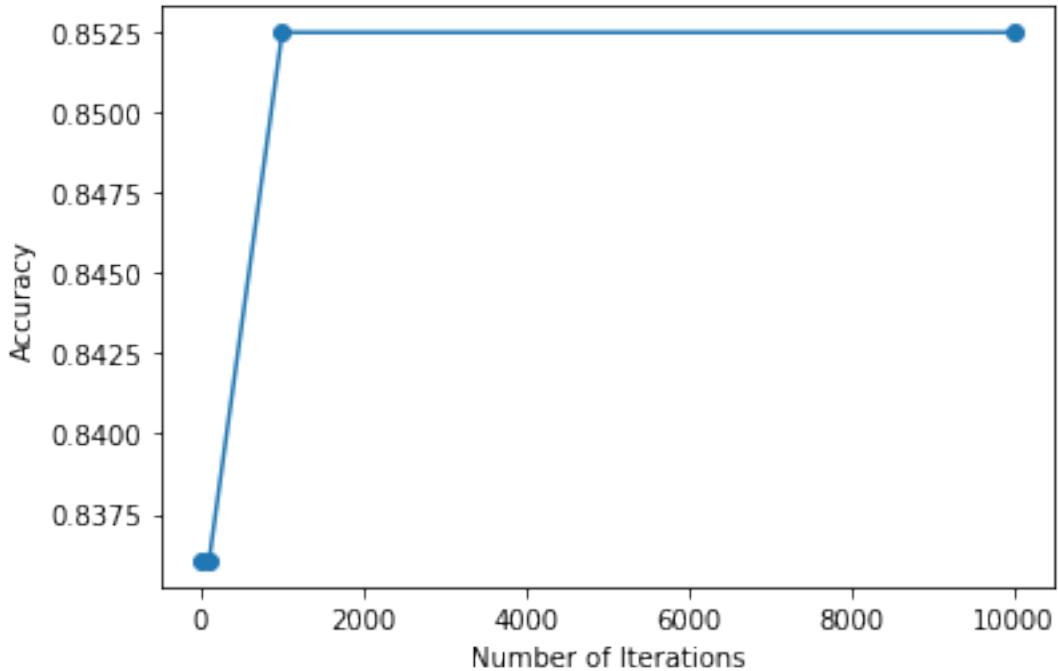
## 8 Number of iterations

```
[18]: its = [10, 100, 1000, 10000]
acc = list()
for it in its:
    model = LogiReg(best_lr, it, False)
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    acc.append(accuracy(y_test, predictions))

best_it = its[acc.index(max(acc))]
print('Best number of iterations is', best_it)
```

Best number of iterations is 1000

```
[19]: plt.plot(its, acc, '-o')
plt.xlabel('Number of Iterations')
plt.ylabel('Accuracy')
plt.savefig('its.jpg')
plt.show()
```



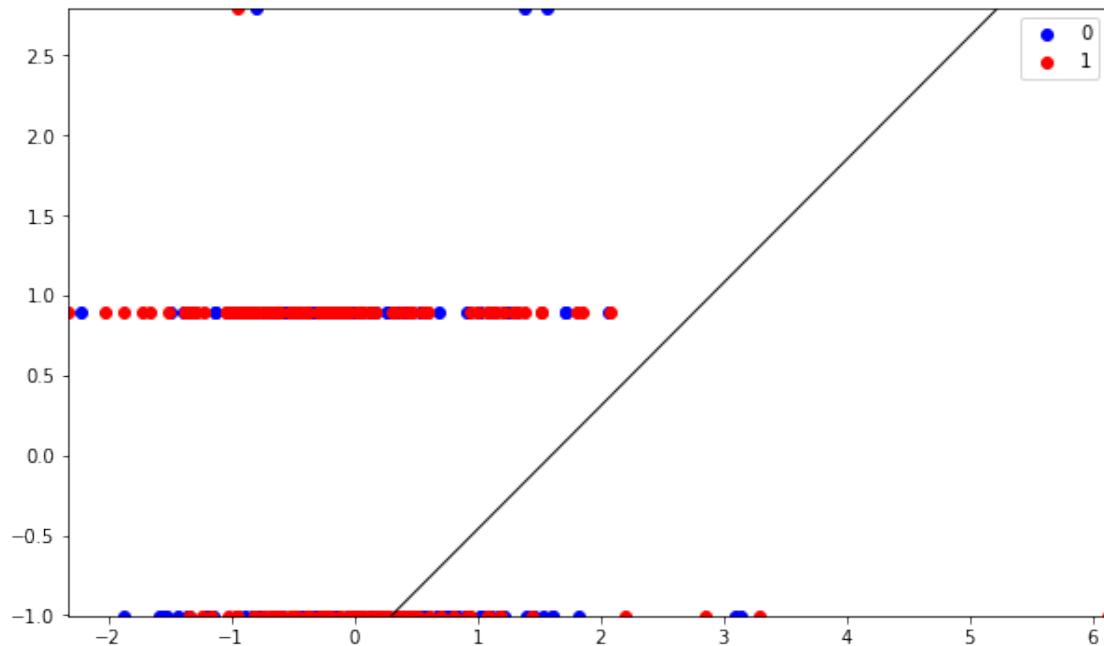
## 9 Feature pairs (random)

```
[20]: def pipeline(X, y, lr=0.01, it=3000):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
    model = LogiReg(lr, it, False)
    model.fit(X_train, y_train)
    return model, round(accuracy(y_test, model.predict(X_test)), 5)*100
```

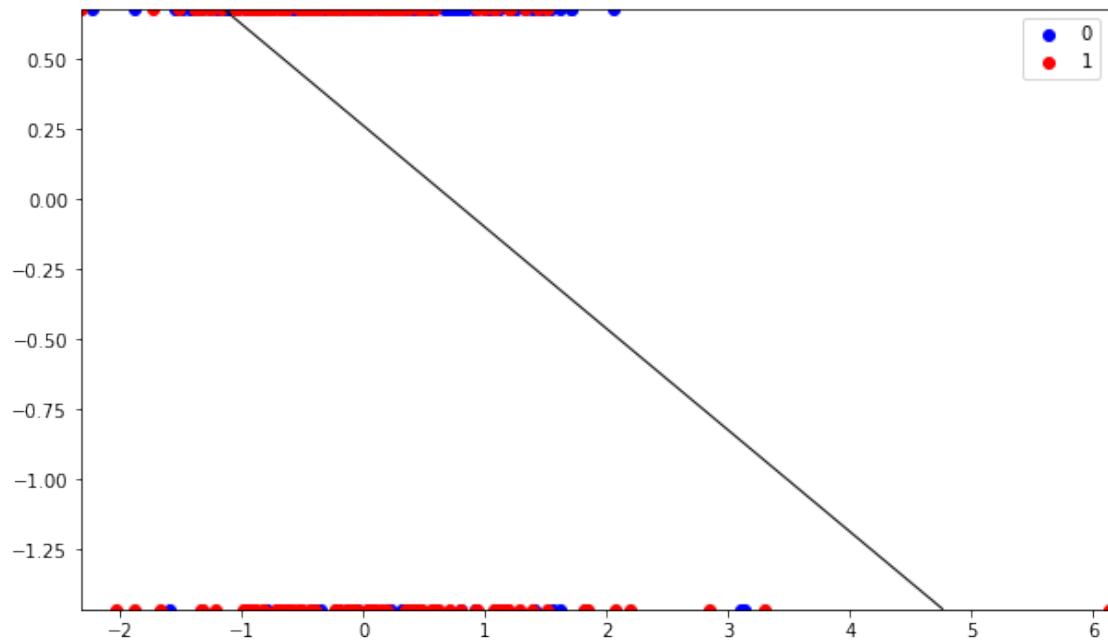
```
[21]: for i in range(5):
    # Pick two features
    f1 = np.random.choice(X.columns.values)
    f2 = f1
    while f2 == f1:
        f2 = np.random.choice(X.columns.values)
    # Train model
    model, acc = pipeline(X[[f1, f2]], y)
```

```
print(f"Trial #{i+1} with features: {f1} and {f2}, accuracy: {acc} %")  
  
newX = X[[f1, f2]].to_numpy()  
plot_regression(newX, y, model)
```

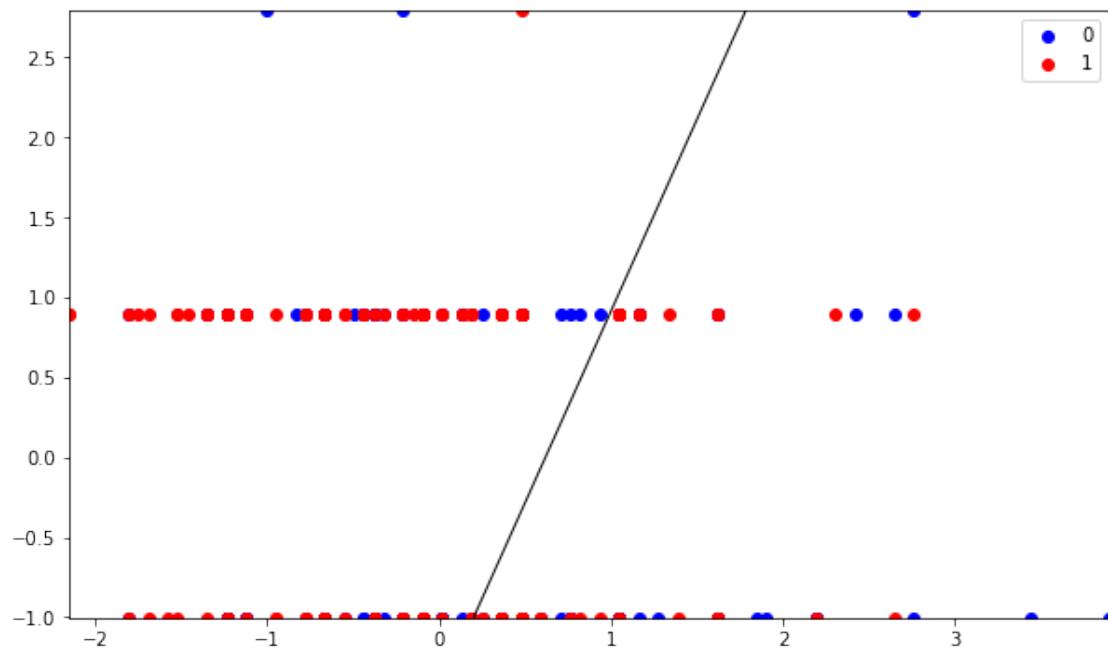
Trial #1 with features: chol and restecg, accuracy: 63.93400000000005 %



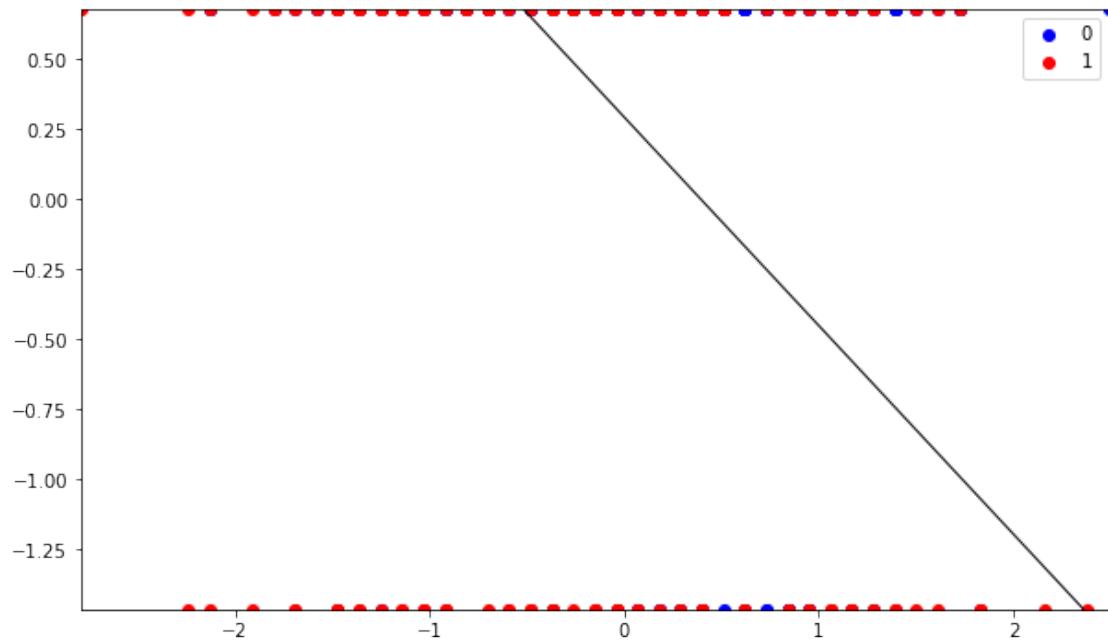
Trial #2 with features: chol and sex, accuracy: 49.18 %



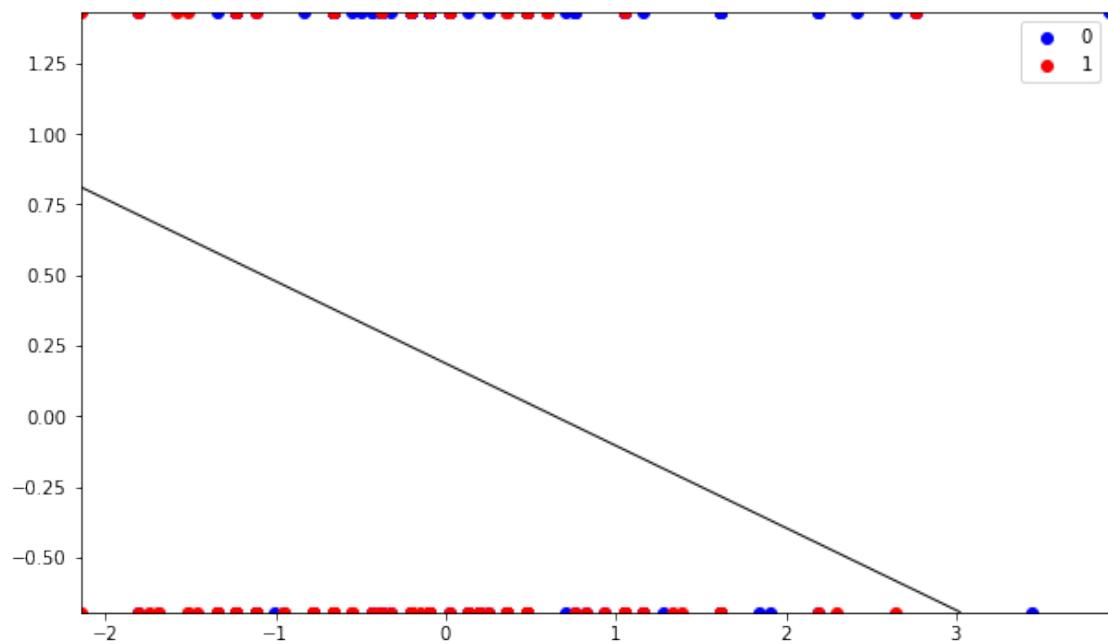
Trial #3 with features: trestbps and restecg, accuracy: 59.016000000000005 %



Trial #4 with features: age and sex, accuracy: 57.377 %



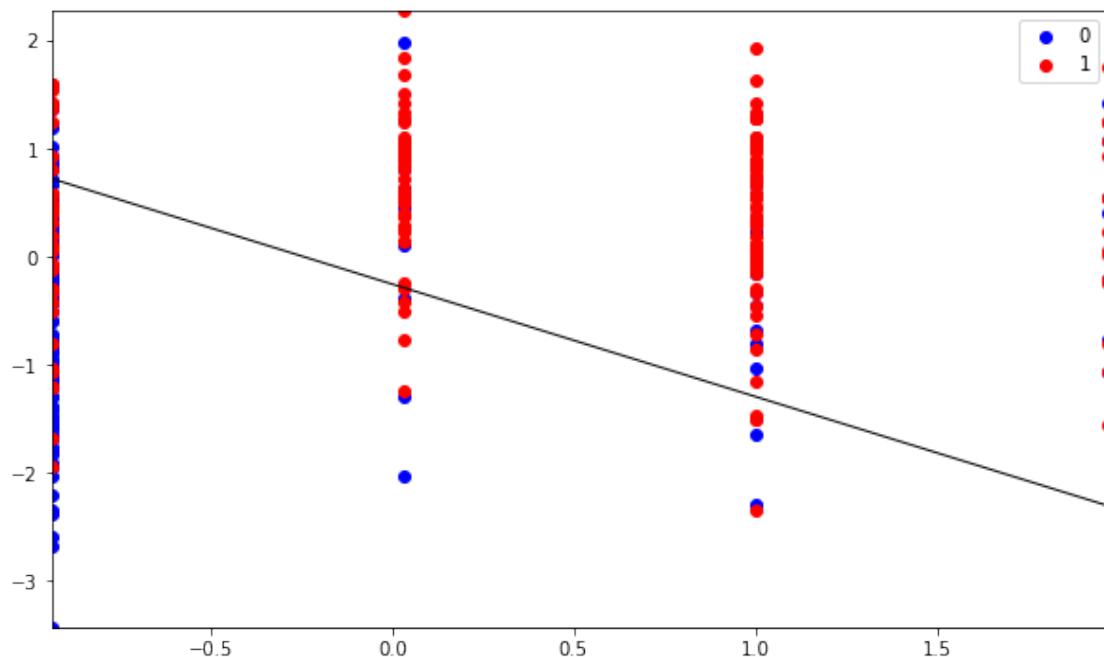
Trial #5 with features: `trestbps` and `exang`, accuracy: 72.131 %



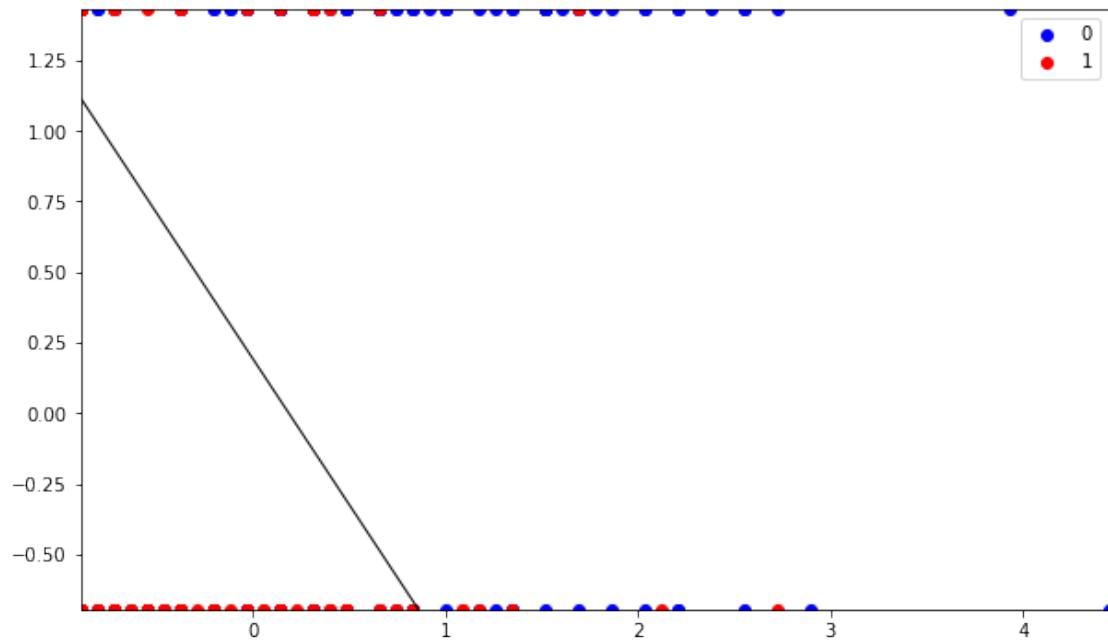
## 10 Feature pairs (based on correlation)

```
[22]: pairs = [['cp', 'thalach'], ['oldpeak', 'exang'], ['ca', 'thal']]  
for i, pair in enumerate(pairs):  
    f1, f2 = pair  
    model, acc = pipeline(X[[f1, f2]], y)  
    print(f"Trial #{i+1} with features: {f1} and {f2}, accuracy: {acc} %")  
  
    newX = X[[f1, f2]].to_numpy()  
    plot_regression(newX, y, model)
```

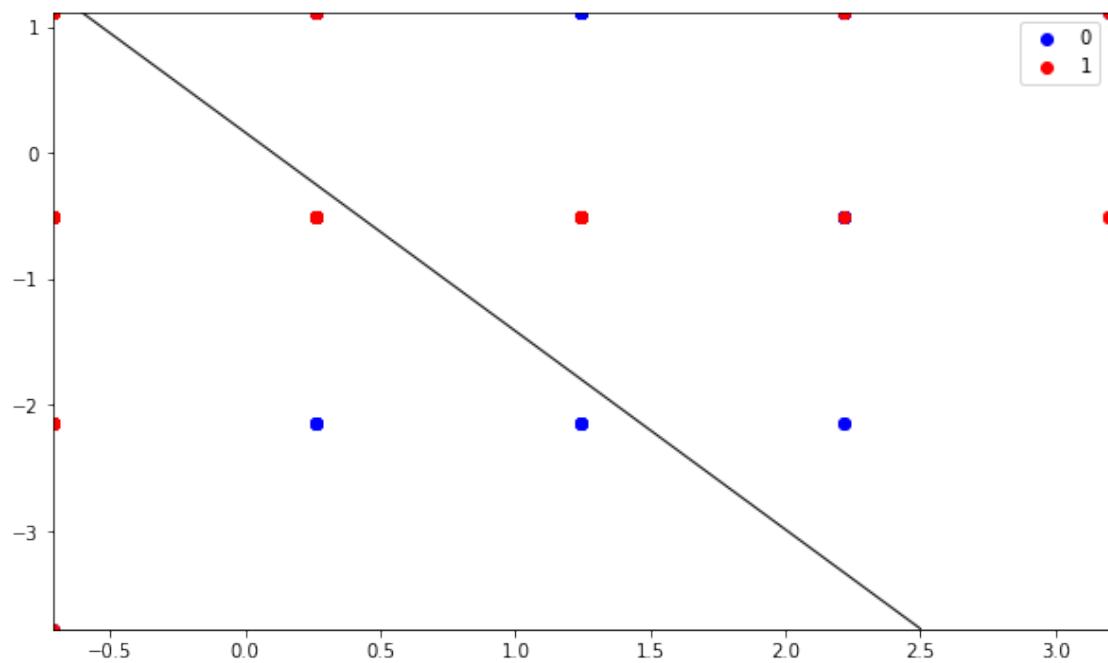
Trial #1 with features: cp and thalach, accuracy: 78.689 %



Trial #2 with features: oldpeak and exang, accuracy: 77.049 %



Trial #3 with features: ca and thal, accuracy: 70.492 %



[ ]:

[ ]:

[ ]:

[ ]:

[ ]: